



# A Characterisation of Convolutional Spiking Autoencoders


Luke Alderson - 02523332



# The Foundations

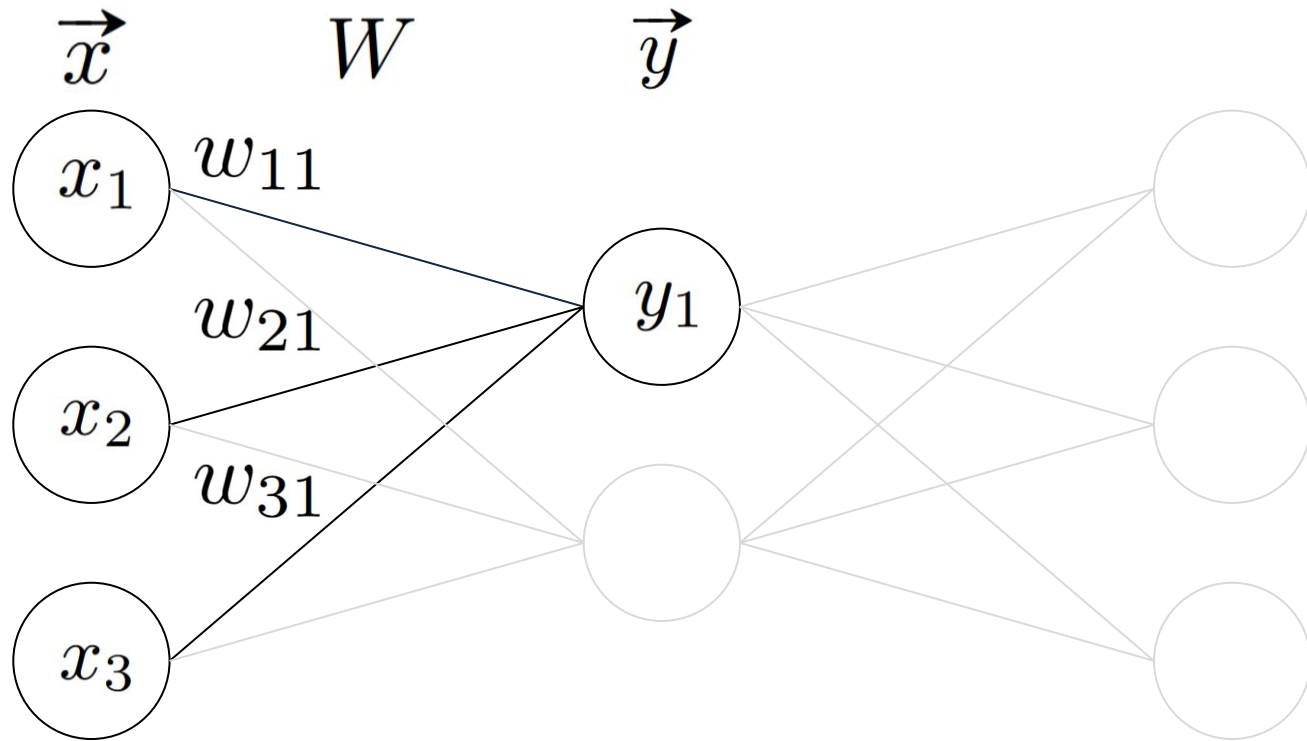


Artificial Neural Network (ANN)




$$\vec{y} = \sigma(W\vec{x} + \vec{b})$$

# The Foundations

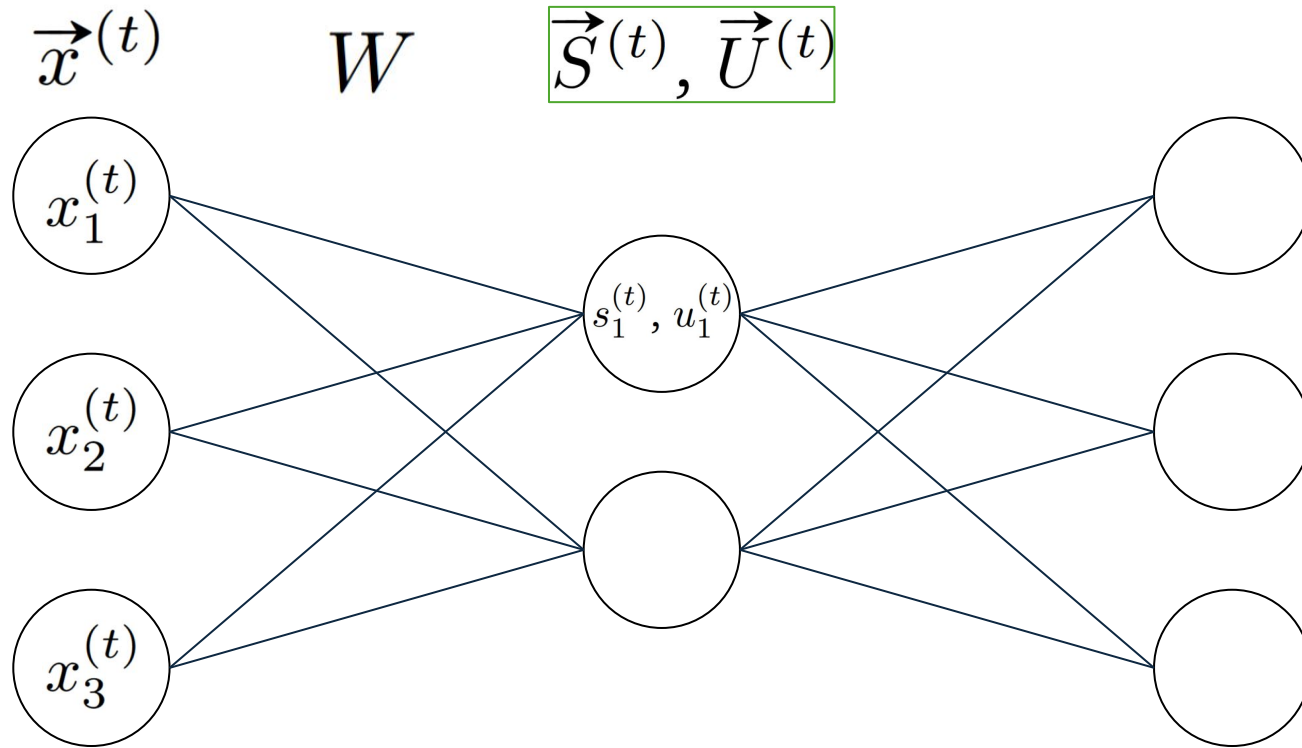


Artificial Neural Network (ANN)



$$\vec{y} = \sigma(W\vec{x} + \vec{b})$$

# Spikes and Surrogate Gradients



Spiking Neural Network (SNN)

$$\tau_m \frac{dU(t)}{dt} = -[U(t) - U_{rest}] + R_m I_{ext}(t)$$


$$U[t] = \beta U[t-1] + (1 - \beta) I_{ext}$$

$$S_{out}(t) = \begin{cases} 1, & U(t) \geq U_{th} \\ 0, & \text{otherwise} \end{cases}$$

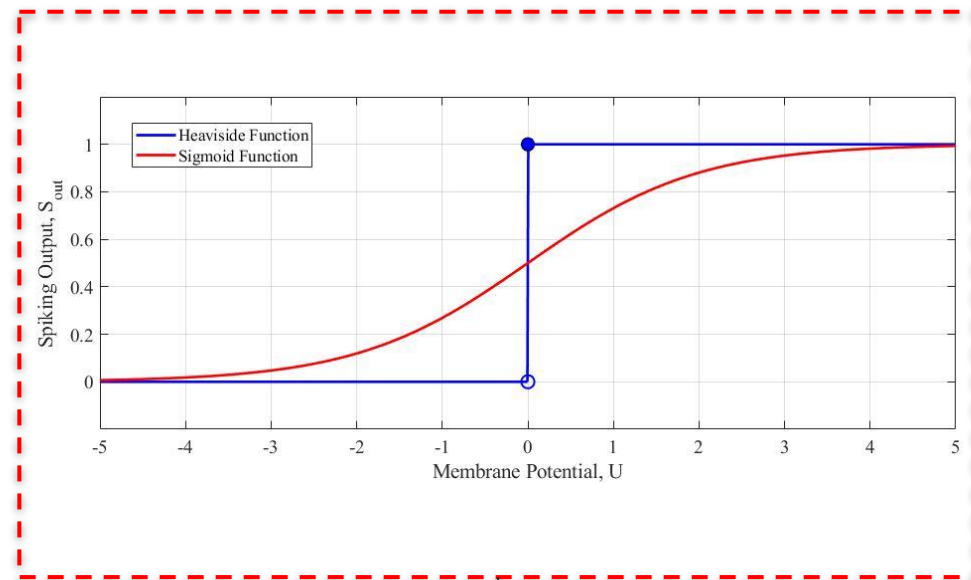
if  $U(t) \geq U_{th}$ , then  $U(t) \rightarrow U_{reset}$

$$\vec{S}^{(t)} = \sigma(\vec{U}^{(t)})$$

# Spikes and Surrogate Gradients




Spiking Neural Network (SNN)



$$\vec{S}^{(t)} = \sigma(\vec{U}^{(t)})$$

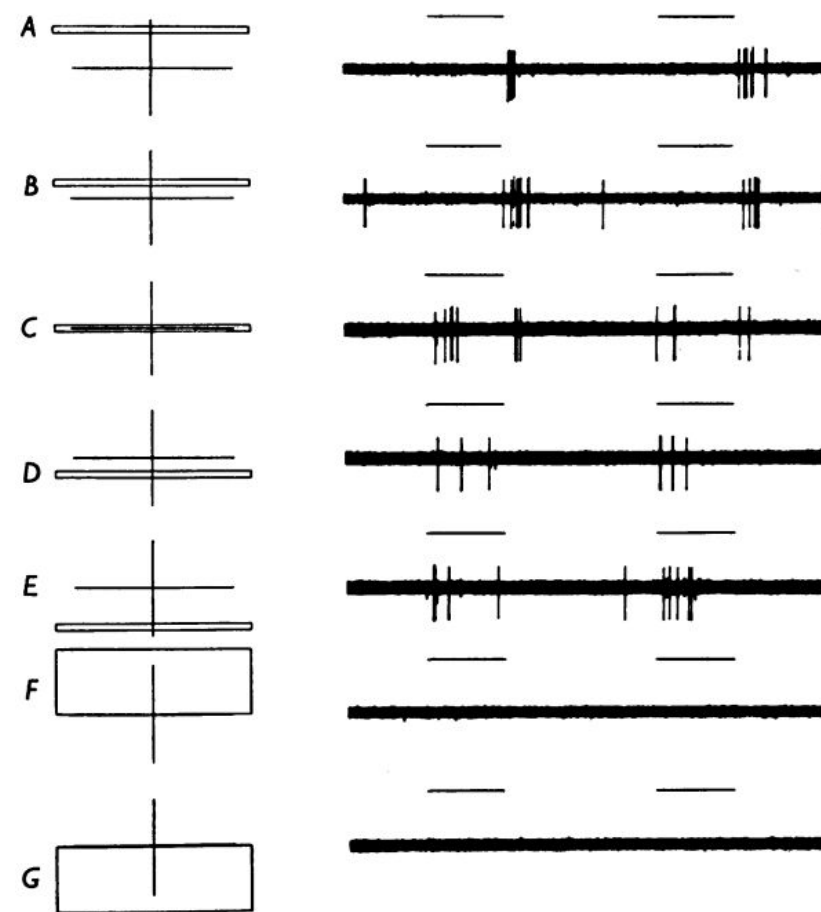
# Artificial and Natural Receptive Fields

## Convolutional Neural Network




*Note: Darker pixels indicate higher activations.*


## Cat Visual Cortex



# Autoencoders




# Convolutional, Spiking Autoencoder





# MNIST - From Floats to Spikes


0 1 2 3 4 5 6 7 8 9



Original



t = 10





t = 50



t = 80



Time Averaged



0 1 2 3 4 5 6 7 8 9

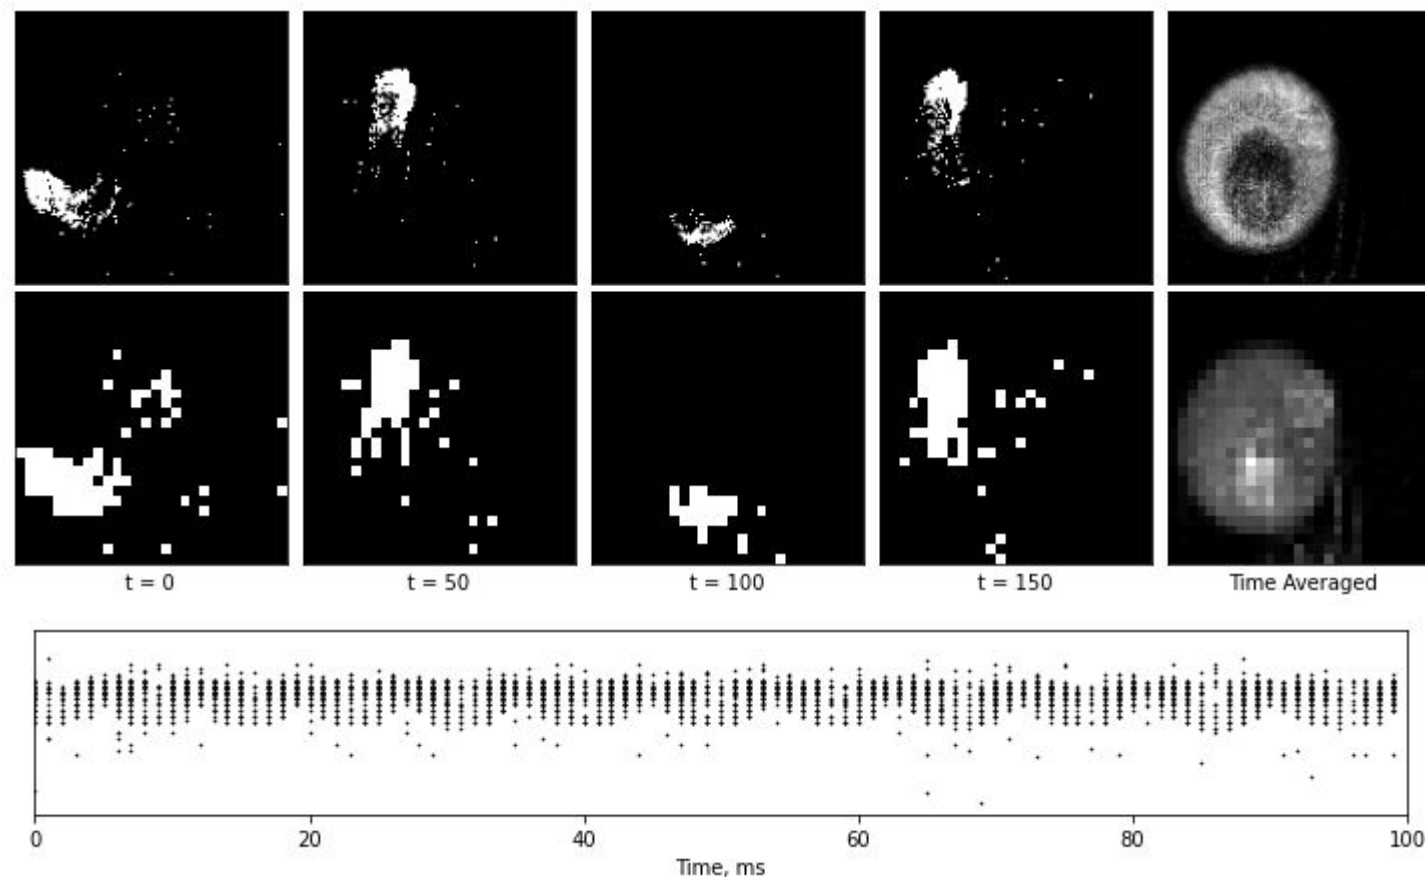

# DVS Gesture - Natively Neuromorphic

Native:

- 128 x 128
- Microsecond


Downsampled:

- 28 x 28
- Millisecond



# Experiments


## Optimal Hyperparameters in SAE and AE




## Encoding Frequency



$$s_i^{(t)} \sim \text{Bin}(T, p_i)$$



## Recurrent Connections




Implicit  
(1-Layer)

Implicit  
(2-Layer)

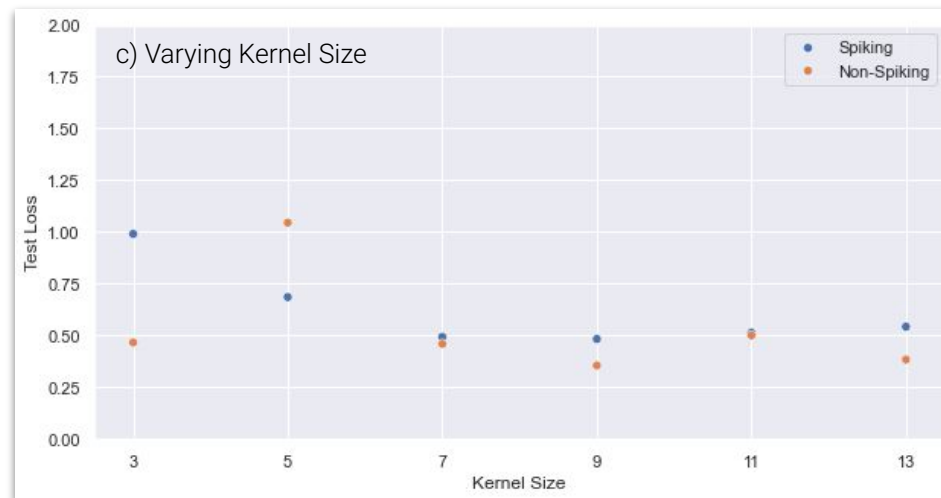
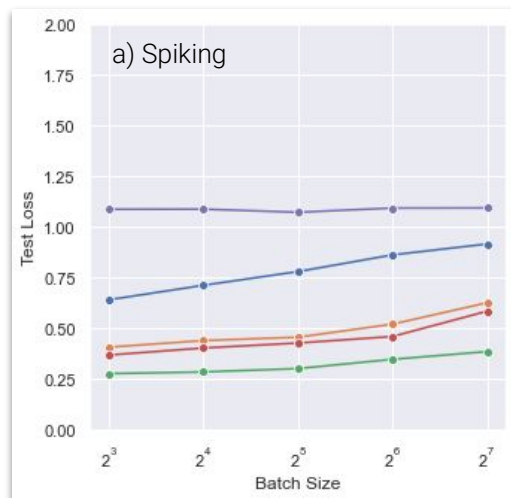
Explicit  
(1-Layer)

## Embedding Size



$$\{x \in \mathbb{N} \mid 2 \leq x \leq 1000\}$$

# Increasing kernel size reduces loss

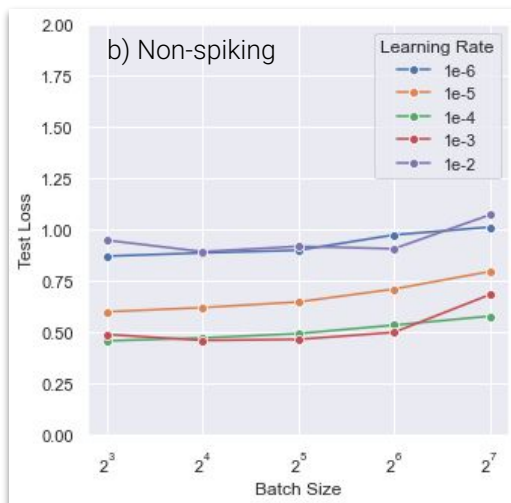


## Optimal Parameters

Batch Size 64

Learning Rate  $1e-4$

Kernel Size 7




d) Original data and reconstructions




# Disparity in rate coding acts as contrast

Spiking Reconstructions

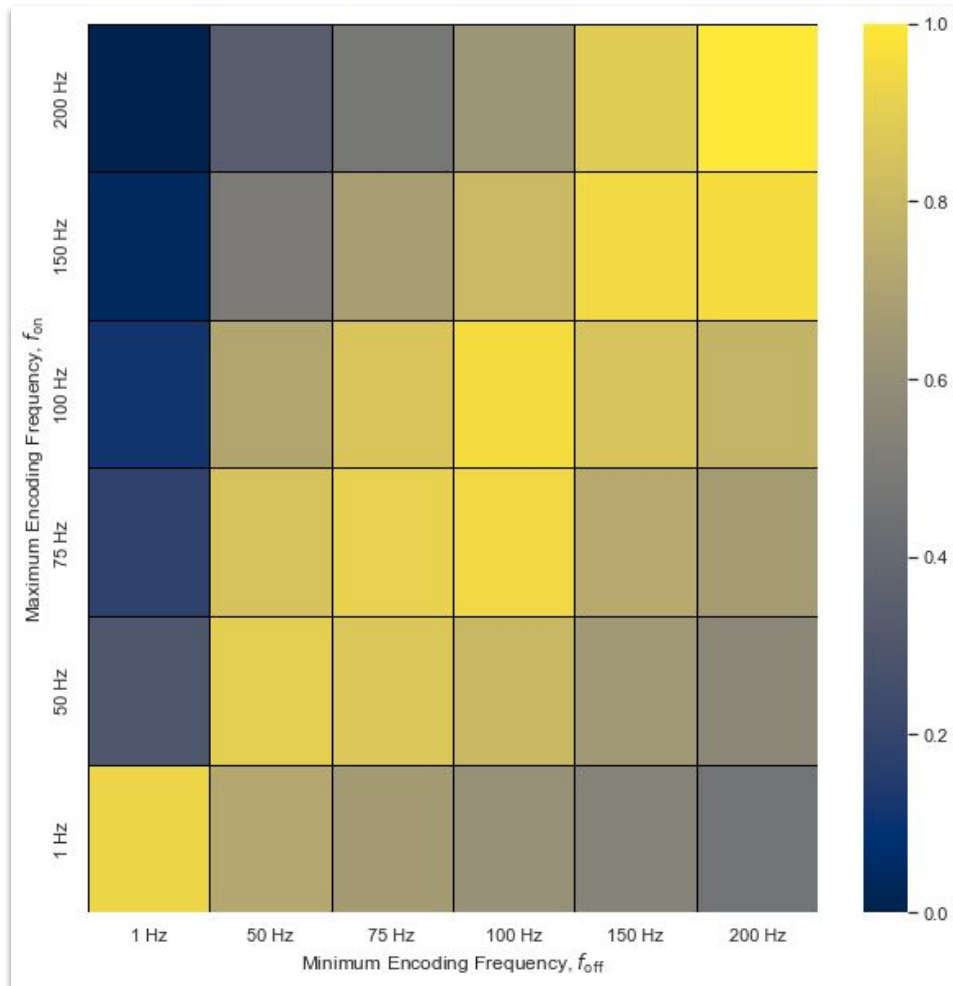


Loss v. Max Encoding Frequency

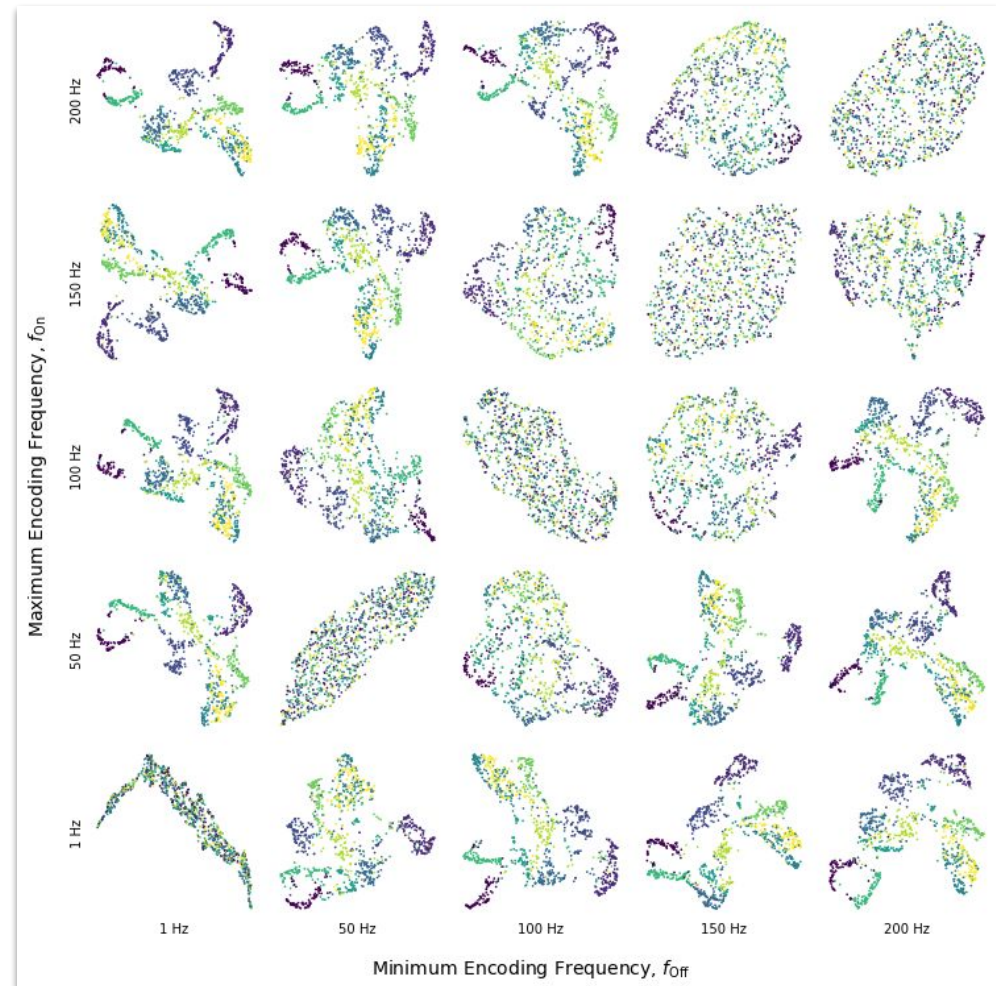


# Disparity in rate coding acts as contrast

Loss Heatmap





UMAP







# Recurrence did not reduce loss



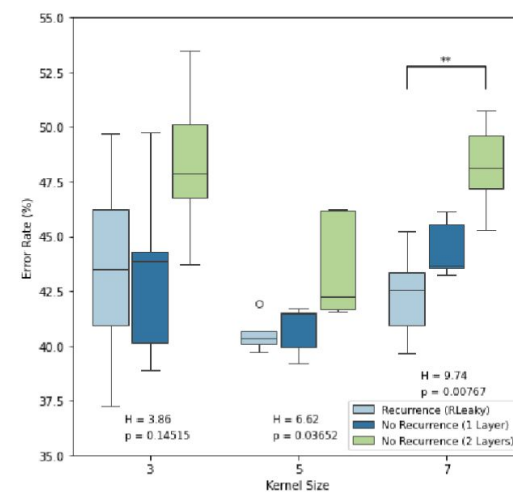
(a) SAE trained on MNIST data with varying kernel size



(b) SAE trained on DVS data with varying kernel size




(c) SAE trained on DVS data with varying training targets




(d) Classifier trained on DVS data with varying kernel size

# Embedding Size generally lowers loss

Spiking Reconstructions

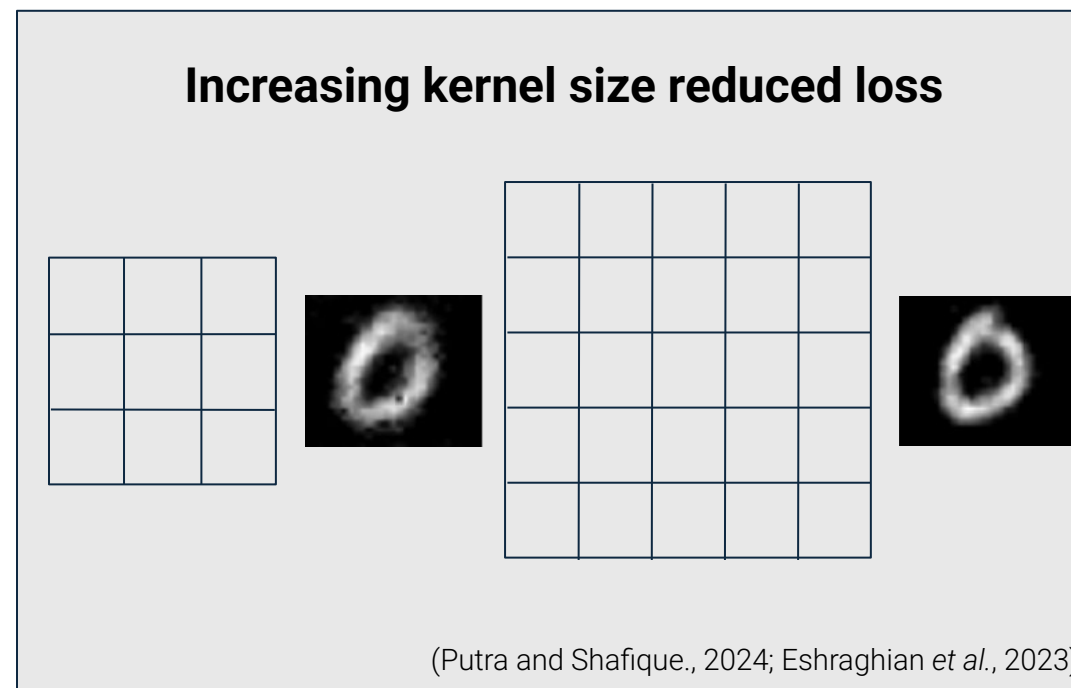



Loss v. Embedding Size



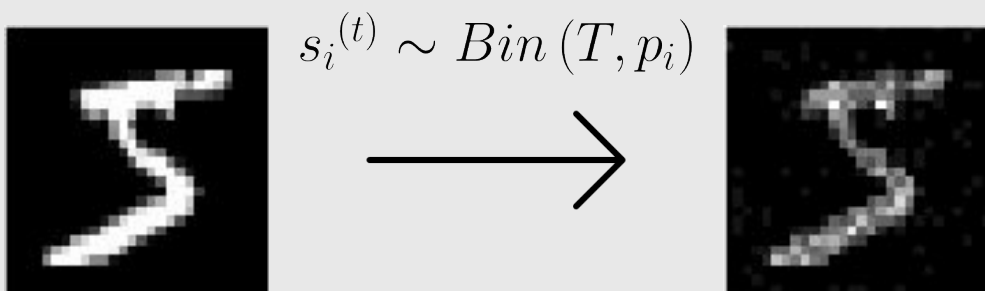


# Discussion - Hyperparameter Sweep





# Discussion - Encoding Strategy

## Encoding Frequency




## Frequency disparity improved embedding quality



(Ben-Shaul et al., 2023; Eshraghian et al., 2023)

# Discussion - Recurrence

## Recurrent Connections



## MNIST


No effect (Bouanane *et al.*, 2023)

## DVS

- 1-Layer, implicit or explicit - no effect.
- Try different loss function.
  - (Cramer *et al.*, 2020)
- 2-Layer - increased loss.
- Vanishing gradients?
  - (Zheng *et al.*, 2021; Ledinauskas *et al.*, 2020)

# Discussion - Embedding Size

## Embedding Size



## Embedding size generally reduced loss

... with difficulty at specific sizes.



Occurred for spiking and non-spiking nets.


(Benmeziane *et al.*, 2023; Dobson, 2023; Li *et al.*, 2017)

# Limitations and Future work

- Only encoding rate was varied.
  - Constrain hidden layer firing rates with regularisation (Hübötter *et al.*, 2021).
- Vanishing gradient problems.
  - E-prop instead of conventional backpropagation (Hoyer *et al.*, 2022).
- Reliance on rate coding.
  - Temporal / Delta / Novel Coding Schemes (Mehta *et al.*, 2002; Ainsworth *et al.*, 2012).

# Thank you for listening!

Any questions?



# Appendices – References

- Ainsworth, M., Lee, S., Cunningham, M.O., Traub, R.D., Kopell, N.J. and Whittington, Miles A. (2012). Rates and Rhythms: A Synergistic View of Frequency and Temporal Coding in Neuronal Networks. *Neuron*, 75(4), pp.572–583. doi:<https://doi.org/10.1016/j.neuron.2012.08.004>.
- Benmeziiane, H., Ounnoughene, Amine Ziad, Hamzaoui, I. and Bouhadjar, Y. (2023). Skip Connections in Spiking Neural Networks: An Analysis of Their Effect on Network Training. *arXiv (Cornell University)*. doi:<https://doi.org/10.48550/arxiv.2303.13563>.
- Ben-Shaul, I., Shwartz-Ziv, R., Galanti, T., Dekel, S. and LeCun, Y. (2023). Reverse Engineering Self-Supervised Learning. *arXiv (Cornell University)*. doi:<https://doi.org/10.48550/arxiv.2305.15614>.
- Burkitt, A.N. (2006). A Review of the Integrate-and-fire Neuron Model: I. Homogeneous Synaptic Input. *Biological Cybernetics*, 95(1), pp.1–19. doi:<https://doi.org/10.1007/s00422-006-0068-6>
- Cramer, B., Yannik Stradmann, Schemmel, J. and Zenke, F. (2022). The Heidelberg Spiking Data Sets for the Systematic Evaluation of Spiking Neural Networks. *IEEE transactions on neural networks and learning systems*, 33(7), pp.2744–2757. doi:<https://doi.org/10.1109/tnnls.2020.3044364>.
- Dobson, J.E. (2023). On reading and interpreting black box deep neural networks. *International Journal of Digital Humanities*, 5(2-3), pp.431–449. doi:<https://doi.org/10.1007/s42803-023-00075-w>.
- Eimantas Ledinauskas, Ruseckas, J., Alfonsas Jursenas and Giedrius Burachas (2020). Training Deep Spiking Neural Networks. *arXiv (Cornell University)*. doi:<https://doi.org/10.48550/arxiv.2006.04436>.
- Eshraghian, J.K., Ward, M., Emre Neftci, Wang, X., Lenz, G., Dwivedi, G., Bennamoun, M., Doo Seok Jeong and Lü, W. (2023). Training Spiking Neural Networks Using Lessons From Deep Learning. *Proceedings of the IEEE*, 111(9), pp.1016–1054. doi:<https://doi.org/10.1109/jproc.2023.3308088>.
- Hoyer, M., Shahram Eivazi and Otte, S. (2022). Efficient LSTM Training with Eligibility Traces. *Lecture notes in computer science*, pp.334–346. doi:[https://doi.org/10.1007/978-3-031-15934-3\\_28](https://doi.org/10.1007/978-3-031-15934-3_28).
- Hubel, D.H. and Wiesel, T.N. (1962). Receptive fields, Binocular Interaction and Functional Architecture in the cat's Visual Cortex. *The Journal of Physiology*, [online] 160(1), pp.106–154. doi:<https://doi.org/10.1113/jphysiol.1962.sp006837>.
- Hübottter, J.F., Lanillos, P. and Tomczak, J.M. (2021). Training Deep Spiking Auto-encoders without Bursting or Dying Neurons through Regularization. *arXiv (Cornell University)*. doi:<https://doi.org/10.48550/arxiv.2109.11045>.
- LeCun, Y., Boser, B., Denker, J.S., Henderson, D., Howard, R.E., Hubbard, W. and Jackel, L.D. (1989). Backpropagation Applied to Handwritten Zip Code Recognition. *Neural Computation*, 1(4), pp.541–551. doi:<https://doi.org/10.1162/neco.1989.1.4.541>.
- Li, H., Xu, Z., Taylor, G., Studer, C. and Goldstein, T. (2018). *Visualizing the Loss Landscape of Neural Nets*. [online] arXiv.org. doi:<https://doi.org/10.48550/arXiv.1712.09913>.
- McInnes, L., Healy, J., Saul, N. and Großberger, L. (2018). UMAP: Uniform Manifold Approximation and Projection. *Journal of Open Source Software*, 3(29), p.861. doi:<https://doi.org/10.21105/joss.00861>.
- Mehta, M.R., Lee, A.K. and Wilson, M.A. (2002). Role of experience and oscillations in transforming a rate code into a temporal code. *Nature*, [online] 417(6890), pp.741–746. doi:<https://doi.org/10.1038/nature00807>.
- Mohamed Sadek Bouanane, Cherifi, D., Chicca, E. and Lyes Khacef (2023). Impact of spiking neurons leakages and network recurrences on event-based spatio-temporal pattern recognition. *Frontiers in Neuroscience*, 17(17). doi:<https://doi.org/10.3389/fnins.2023.1244675>.
- Putra, and Shafique, M. (2024). A Methodology for Improving Accuracy of Embedded Spiking Neural Networks through Kernel Size Scaling. *arXiv (Cornell University)*. doi:<https://doi.org/10.48550/arxiv.2404.01685>.
- Zheng, H., Wu, Y., Deng, L., Hu, Y. and Li, G. (2021). Going Deeper With Directly-Trained Larger Spiking Neural Networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(12), pp.11062–11070. doi:<https://doi.org/10.1609/aaai.v35i12.17320>.

# Appendices - CNN

$X$

1	2	3	4	2	1
4	5	6	5	0	1
7	8	9	2	2	1
9	3	2	4	2	1
5	7	3	2	1	9
4	4	3	2	8	1

$W$

1	1	0
0	0	0
0	0	1

$+ b$

$Y$

12			

$$Y = \sigma(W \cdot X + b)$$

Assume  $b = 0$  and the activation function is the identity.



# Appendices - CNN

$X$

1	2	3	4	2	1
4	5	6	5	0	1
7	8	9	2	2	1
9	3	2	4	2	1
5	7	3	2	1	9
4	4	3	2	8	1

$W$

1	1	0
0	0	0
0	0	1

$+ b$

$Y$

12	7		

$$Y = \sigma(W \cdot X + b)$$

Assume  $b = 0$  and the activation function is the identity.

# Appendices - CNN

 $X$ 

1	2	3	4	2	1
4	5	6	5	0	1
7	8	9	2	2	1
9	3	2	4	2	1
5	7	3	2	1	9
4	4	3	2	8	1

 $W$ 

1	1	0
0	0	0
0	0	1

 $+ b$  $Y$ 

12	7	9	7
11	15	13	6
18	19	12	13
15	7	14	7

$$Y = \sigma(W \cdot X + b)$$

Assume  $b = 0$  and the activation function is the identity.

# Appendices - Loss Function

$$L(\theta)_{nRMSE} = \frac{\sqrt{\frac{1}{N} \sum_{i=1}^N (X_i - \hat{X}_i)^2}}{\sqrt{\frac{1}{N} \sum_{i=1}^N (X_i - \bar{X})^2}} = \frac{L(\theta)_{RMSE}}{\sigma_X}$$

where

$L(\theta)_{RMSE}$  is the RMSE loss as a function of the network parameters.

$N$  is the number of neurons.

$X_i$  is the input spike count of the  $i$ th neuron.

$\hat{X}_i$  is the output spike count of the  $i$ th neuron.

$\bar{X}$  is the mean input spike count.

$\sigma_X$  is the standard deviation of the input data.

# Appendices – Spike Encoding

$$p_i = \Delta t[(f_{on} - f_{off})x_i + f_{off}] \quad (6)$$

where

$\Delta t$  is the size of the time-step for the simulation

$x_i$  is the normalised value of the  $i$ th pixel in the input image

$f_{off} \in [0, \frac{1}{\Delta t}]$  is the minimum firing rate in Hz

$f_{on} \in [0, \frac{1}{\Delta t}]$  is the maximum firing rate in Hz

$$S_i^{(t)} \sim \text{Bin}(T, p_i) \quad (5)$$

where

$S_i^{(t)} \in \{0, 1\}$  is the random variable describing spike generation from neuron  $i$  at time  $t$

$T$  is the total number of time-steps in the spike train

$p_i$  is the probability of a spike occurring from neuron  $i$

# Appendices – Backpropagation

Backpropagation faces challenges when applied directly to SNNs because spikes are modelled as instances of the Dirac-delta function, making them non-differentiable. The general form of the chain rule for a SNN is given by:

$$\frac{\partial L}{\partial W_{ij}} = \frac{\partial L}{\partial S_{out}} \frac{\partial S_{out}}{\partial U} \frac{\partial U}{\partial W_{ij}} \quad (2)$$

where  $L$  is the loss function, such as the Euclidean norm,  $W_{ij}$  represents the weight between the  $i$ th neuron of a given layer and the  $j$ th neuron in the subsequent layer, and  $S_{out}$  is the spike released from the neuron as a function of the membrane potential,  $U$ .

The weight update is then computed as:

$$W_{ij}^{t+1} = W_{ij}^t - \alpha \frac{\partial L}{\partial W_{ij}^t} \quad (3)$$

where  $W_{ij}^{t+1}$  is the updated weight,  $W_{ij}^t$  is the current weight, and  $\alpha$  is the learning rate.

Direct application of the chain rule in spiking neuron layers is problematic because it ultimately prevents useful weight updates as described below:

$$\frac{\partial S_{out}}{\partial U} \in \{0, \infty\} \implies \frac{\partial L}{\partial W_{ij}^t} \in \{0, \infty\} \implies W_{ij}^{t+1} \in \{W_{ij}^t, -\infty\} \quad (4)$$